
Обучение модели с бинарными весами для задачи классификации голосовых команд.

Задачи


- Подготовка датасета - конвертация команд с FFT, мы будем работать в этом домене, это нам позволит использовать признаки фиксированного размера.
- Обучение модели для классификации с бинарными весами.
- Написание кернелей на C++ под CPU для матричных умножений.
- Транспортировка полученных весов в C++ кернели
- Сравнение скорости и качества полученной модели с моделью на pytorch.

Mel spectrograms



паддинг

`torchaudio.transforms.
MelSpectrogram`

`torchaudio.transforms.
AmplitudeToDB`



Обучение маленькой сверточной сети без бинаризации



XNOR-Net: ImageNet Classification Using Binary Convolutional Neural Networks

Estimating binary weights: Without loss of generality we assume \mathbf{W}, \mathbf{B} are vectors in \mathbb{R}^n , where $n = c \times w \times h$. To find an optimal estimation for $\mathbf{W} \approx \alpha\mathbf{B}$, we solve the following optimization:

$$\begin{aligned} J(\mathbf{B}, \alpha) &= \|\mathbf{W} - \alpha\mathbf{B}\|^2 \\ \alpha^*, \mathbf{B}^* &= \underset{\alpha, \mathbf{B}}{\operatorname{argmin}} J(\mathbf{B}, \alpha) \end{aligned} \tag{2}$$

by expanding equation 2, we have

$$J(\mathbf{B}, \alpha) = \alpha^2 \mathbf{B}^\top \mathbf{B} - 2\alpha \mathbf{W}^\top \mathbf{B} + \mathbf{W}^\top \mathbf{W} \tag{3}$$

since $\mathbf{B} \in \{+1, -1\}^n$, $\mathbf{B}^\top \mathbf{B} = n$ is a constant. $\mathbf{W}^\top \mathbf{W}$ is also a constant because \mathbf{W} is a known variable. Lets define $\mathbf{c} = \mathbf{W}^\top \mathbf{W}$. Now, we can rewrite the equation 3 as follow: $J(\mathbf{B}, \alpha) = \alpha^2 n - 2\alpha \mathbf{W}^\top \mathbf{B} + \mathbf{c}$. The optimal solution for \mathbf{B} can be achieved by maximizing the following constrained optimization: (note that α is a positive value in equation 2, therefore it can be ignored in the maximization)

$$\mathbf{B}^* = \underset{\mathbf{B}}{\operatorname{argmax}} \{ \mathbf{W}^\top \mathbf{B} \} \quad s.t. \quad \mathbf{B} \in \{+1, -1\}^n \tag{4}$$


This optimization can be solved by assigning $\mathbf{B}_i = +1$ if $\mathbf{W}_i \geq 0$ and $\mathbf{B}_i = -1$ if $\mathbf{W}_i < 0$, therefore the optimal solution is $\mathbf{B}^* = \operatorname{sign}(\mathbf{W})$. In order to find the optimal value for the scaling factor α^* , we take the derivative of J with respect to α and set it to zero:

$$\alpha^* = \frac{\mathbf{W}^\top \mathbf{B}^*}{n} \tag{5}$$

By replacing \mathbf{B}^* with $\operatorname{sign}(\mathbf{W})$

$$\alpha^* = \frac{\mathbf{W}^\top \operatorname{sign}(\mathbf{W})}{n} = \frac{\sum |\mathbf{W}_i|}{n} = \boxed{\frac{1}{n} \|\mathbf{W}\|_{\ell_1}} \tag{6}$$

XNOR-Net: ImageNet Classification Using Binary Convolutional Neural Networks



XNOR-Net: ImageNet Classification Using Binary Convolutional Neural Networks

- gradient for sign function
- порядок слоев

$$\frac{\partial \text{sign}}{\partial r} = r 1_{|r| \leq 1}$$
$$\frac{\partial C}{\partial W_i} = \frac{\partial C}{\widetilde{W}_i} \left(\frac{1}{n} + \frac{\partial \text{sign}}{\partial W_i} \alpha \right)$$





Fig. 3: This figure contrasts the block structure in our XNOR-Network (right) with a typical CNN (left).

Наивная реализация XNOR-Net

```
class SignSTE(torch.autograd.Function):
    @staticmethod
    def forward(ctx, x):
        ctx.save_for_backward(x)
        return x.sign()
    @staticmethod
    def backward(ctx, g):
        (x,) = ctx.saved_tensors
        mask = (x.abs() <= 1).to(g.dtype)
        return g * mask
```

```
class NaiveXNORLinear(nn.Module):
    def __init__(self, ch_in, ch_out):
        super(NaiveXNORLinear, self).__init__()
        self.fc = nn.Linear(ch_in, ch_out, bias=True)



    def forward(self, input_x):
        quantized_weight = SignSTE.apply(self.fc.weight)
        quintized_input = SignSTE.apply(input_x)
        out = nn.functional.linear(quintized_input, quantized_weight)
        alpha = self.fc.weight.abs().mean()
        betta = input_x.abs().mean()
        return out*alpha*betta + self.fc.bias

class NaiveXNORConv2d(nn.Module):
    def __init__(self, ch_in, ch_out, kernel=3, padding=0, stride=1):
        super(NaiveXNORConv2d, self).__init__()
        self.padding = padding
        self.stride = stride
        self.conv = nn.Conv2d(ch_in, ch_out, kernel, stride, padding, bias=False)
        k = torch.ones((1, 1, kernel, kernel))/(kernel**2) #< Now consider only constants
        self.k = nn.Parameter.Buffer(k, persistent=False)

    def forward(self, input_x):
        # A, k - BinActiv
        A = input_x.abs().mean(axis=1, keepdim=True)
        K = nn.functional.conv2d(A, self.k, None, self.stride, self.padding)
        alpha = self.conv.weight.abs().mean()

        quantized_weight = SignSTE.apply(self.conv.weight)
        quintized_input = SignSTE.apply(input_x)
        out = nn.functional.conv2d(quintized_input, quantized_weight, None, self.stride, self.padding)
        return out*K*alpha
```

Наивная реализация XNOR-Net



Смотрим готовую реализацию

- <https://github.com/allenai/XNOR-Net> - оригинальный репозиторий на lua
- <https://github.com/jiecaoyu/XNOR-Net-PyTorch>

- class BinOp - бинаризация и градиент
- class BinActive - Sign forward, backward
- class BinConv2d - свертка и линейный слой (2 в одном) + batch norm + relu

```



bin_op = BinOp(model) #< use bin_op

for epoch in range(num_epochs):
    for i, (series, labels) in tqdm.tqdm(enumerate(train_loader), total=len(train_loader)):
        Net.train()
        bin_op.binarization() #< use bin_op
        outputs = Net(series.to(device))
        loss = loss_func(outputs, labels.to(device))
        self.loss_list.append(loss.item())

        optimizer.zero_grad()
        loss.backward()
        bin_op.restore() #< use bin_op
        bin_op.updateBinaryGradWeight() #< use bin_op

        optimizer.step()
    
```

Смотрим готовую реализацию



Небольшой рефакторинг

```
@staticmethod
def forward(ctx, weight: torch.Tensor):
    if weight.dim() == 4:
        # [out, in, kh, kw]
        negMean = weight.mean(dim=1, keepdim=True).mul(-1).expand_as(weight)
    elif weight.dim() == 2:
        # [out, in]
        negMean = weight.mean(dim=1, keepdim=True).mul(-1).expand_as(weight)

    w_centered = weight + negMean
    w_clamped = w_centered.clamp_(-1.0, 1.0)

    if weight.dim() == 4:
        n = weight[0].numel()
        m = (w_clamped.norm(p=1, dim=3, keepdim=True)
              .sum(2, keepdim=True)
              .sum(1, keepdim=True)
              .div(n))
    else:
        n = weight.size(1)
        m = w_clamped.norm(p=1, dim=1, keepdim=True).div(n)

    w_bin = w_clamped.sign()

    ctx.save_for_backward(w_clamped)
    ctx.n = n
    return w_bin, m
```

```
@staticmethod
def backward(ctx, grad_out: torch.Tensor, scale_grad:torch.Tensor):
    (w_clamped,) = ctx.saved_tensors
    n = ctx.n
    s = w_clamped.size()

    g = grad_out
    if w_clamped.dim() == 4:
        m = (w_clamped.norm(p=1, dim=3, keepdim=True)
              .sum(2, keepdim=True)
              .sum(1, keepdim=True)
              .div(n)).expand(s)
    else:
        m = w_clamped.norm(p=1, dim=1, keepdim=True).div(n).expand(s)

    m = m.clone()
    m[w_clamped.lt(-1.0)] = 0
    m[w_clamped.gt(1.0)] = 0

    grad = m * g



    m_add = w_clamped.sign() * g
    if w_clamped.dim() == 4:
        m_add = (m_add.sum(3, keepdim=True)
                  .sum(2, keepdim=True)
                  .sum(1, keepdim=True)
                  .div(n)
                  .expand(s))
    else:
        m_add = m_add.sum(1, keepdim=True).div(n).expand(s)

    m_add = m_add * w_clamped.sign()
    grad = (grad + m_add) * (1.0 - 1.0 / s[1]) * n
    #grad = grad * 1e9
    return grad
```



<https://github.com/allenai/XNOR-Net/issues/41>

<https://github.com/jiecaoyu/XNOR-Net-PyTorch/issues/126>

Использование квантизации из семинара



сравнение матричных умножений




<https://habr.com/ru/articles/359272/>

<https://stackoverflow.com/questions/79806482/why-is-eigen-c-int-matrix-multiplication-10x-slower-than-float-multiplication>

torch.utils.cpp_extension

- несколько подходов: можно сделать через `load_inline`, можно через `setup`
- <https://github.com/pytorch/extension-cpp>
- https://github.com/AmirMardan/pytorch_extending_cpp_binding
- могут возникнуть проблемы, если `setuptools>75.8.2`
<https://github.com/pypa/setuptools/issues/4874>
- флаг "`-march=native`" - немного ускоряет кастомные реализации


XNOR + POPCOUNT



$a_{ij} \in [-1, 1]$

\rightarrow

$P_{ij} - \text{int } 64$





XNOR + POPCOUNT


```
tensor([ 1., -1.,  1.,  1., -1.,  1.,  1., -1., -1.,  1., -1., -1.,
        -1.,  1.,  1.,  1., -1.,  1.,  1., -1., -1.,  1., -1., -1.,
        -1.,  1., -1.,  1.,  1.,  1., -1., -1.,  1.,  1.,  1., -1.,
        -1., -1.,  1., -1.,  1., -1., -1., -1.,  1.,  1., -1., -1.,
        1., -1., -1.,  1., -1., -1., -1.])
tensor(-8571384234711997843)
```

```
100010010000110001010011011001110100000100110111010001001101101
```

$$\sum_{i=1}^K a_i \cdot b_i = (+1) \cdot (equal) + (-1) \cdot (diff) = equal - (K - equal) = 2 \cdot equal - K$$

сравнение с pytorch





Другие реализации

- <https://github.com/honcharov-danylo/ultraspeed>
- <https://github.com/cooooorn/Pytorch-XNOR-Net>
- https://github.com/brycexu/BNN_Kernel
- <https://github.com/tairenpiao/XNOR-popcount-GEMM-PyTorch-CPU-CUDA>